
The Chain of Thought as Compositional Memory in Unmodified Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Unmodified pretrained language models use the hidden states at written
2 intermediate-value tokens as compositional registers for serial task state. Over-
3 writing the hidden state at one written value’s position with the state a different
4 number would have produced, while leaving the visible text unchanged, redirects
5 the final answer to the never-written number in 0.70 to 1.00 of cases across six
6 models from four families on arithmetic and narrative state-tracking tasks; a size-
7 matched random perturbation redirects at most 0.01 for instruction-tuned models.
8 The value is retrieved through attention to its own position: blocking that attention
9 collapses the answer’s dependence on the value to at most 0.007 on four models,
10 matched blocks of neighboring, random, and earlier-value positions change nothing,
11 and the model continues to produce fluent parseable answers. A single-layer
12 version of the patch succeeds at any depth up to roughly two thirds of the network
13 and fails after, locating the retrieval before that point. Two positions function as
14 independent stores: values planted at both compose into their sum, a number appearing
15 nowhere in the input, at 0.88 to 0.90, and answers drawing on only one
16 planted value occur at 0.01. A calibrated ablation of the concepts a fitted Jacobian
17 lens reads from the residual stream (J-space), at the strongest dose preserving ordinary
18 generation, leaves this externalized computation intact, bounding what that
19 readout captures rather than locating serial state outside the workspace it samples.
20 These results carry a caution for oversight: the written positions are a causally
21 used memory channel, and identical visible text can coexist with different stored
22 values and different downstream behavior, so a trace exposes the symbols of the
23 computation without guaranteeing their hidden content.

24 1 Introduction

25 When a pretrained language model tracks serial state through a chain of thought, the hidden state at
26 each written value’s token is a memory register: settable by intervention, retrieved through attention
27 to that position, and independently addressable, with the stored value rather than the visible text
28 determining the answer. We establish this with causal interventions on six unmodified models from
29 four families (Qwen, Phi, OLMo, DeepSeek-R1-Distill), on two task families with exact ground
30 truth for every intermediate value, and we locate the retrieval in the network’s depth.

31 Prior work anticipated pieces of this picture without establishing the mechanism in the wild. Ex-
32 pressivity theory treats generated tokens as the recurrent state of an otherwise fixed-depth compu-
33 tation [Merrill and Sabharwal, 2024], Self-Notes demonstrates writing as working memory behav-
34 iorally [Lanchantin et al., 2023], Zhu et al. [2025] show that intermediate-result tokens store muta-
35 ble values on multiplication and dynamic-programming tasks, and Shih et al. [2026] patch written
36 scratchpad representations on a fine-tuned model. None of these show that unmodified pretrained

models do this natively, identify the retrieval route, test whether written stores compose, or map when a model reads a written value rather than recomputing it. Those four questions are this paper’s results:

1. **The hidden state at a written value’s position is a value register, natively.** Planting a never-written value into the residual stream at that position, text unchanged, redirects the answer to the planted value at 0.70 to 1.00 across all six models; restoring the correct state under corrupted text reverts the answer at 0.88 to 1.00; size-matched random noise redirects at most 0.01 on instruction-tuned models (Section 3).
2. **Retrieval runs through attention to the written position, before two-thirds depth.** Blocking attention from later positions to that one position collapses the answer’s dependence on the value to at most 0.007 on four models in three families, with matched control blocks at the unblocked level and fluency unchanged. Planting through a single layer succeeds at any layer up to two thirds of the network and fails at deeper layers, so the value is consumed before that point (Section 4).
3. **Written positions are independently addressable.** Two never-written values planted at two positions yield their sum, absent from the input, at 0.88 to 0.90; answers using only one of the two occur at 0.01 (Section 3).
4. **A calibrated J-space ablation does not disturb the externalized computation.** Removing the top concepts a fitted Jacobian lens reads from the residual stream, at the strongest dose that preserves ordinary generation, changes neither accuracy nor what gets written. This bounds the particular readout tested; it does not place serial state outside the workspace the lens samples (Section 5).
5. **Reading is the default; recomputation is model-specific.** Models follow an edited written value even when the prompt implies the correct one. DeepSeek V3.2 recomputes only values the prompt generatively defines, and precisely those; Claude Sonnet 4.5 also uses stated checkpoints; Llama 3.3 70B recomputes almost nothing despite solving the same chains from scratch (Section 6).

The picture these results support is specific: on serial state-tracking tasks, pretrained models spontaneously write task state into the trace and retrieve it from the written positions through attention, and the stored values behave as independently settable registers. The visible token and the stored state normally agree, and our interventions show they can be dissociated: the same text can carry different hidden values with different downstream consequences, which qualifies what reading a trace can establish about the computation behind it.

2 Tasks and protocol

Both task families give exact ground truth for every intermediate value, so the downstream effect of any edit is computable. Arithmetic chains start from a three-digit number and apply d signed two-digit additions and subtractions. Narrative state-tracking problems describe a character whose count of an item changes through varied events, with varied names, items, verbs, and sentence forms. In both, the model writes one line per step giving the running value.

The protocol takes a correct worked solution, changes the value at one step j , truncates immediately after the edited line, and lets the model continue (Figure 1). A no-edit truncation measures the sampling noise floor, which is 0.00 throughout. Generation uses temperature 0.6 and top- p 0.95 with five samples per item; answers are parsed from a fixed final-line format and unparseable outputs are counted separately. Reported intervals are 95 percent bootstrap intervals resampling items; we report intervals rather than significance tests, and the effects that carry the argument exceed their controls by margins far larger than the intervals. White-box experiments use Qwen3-4B, Qwen2.5-7B-Instruct, Phi-3-medium, OLMo-2-7B-Instruct, and DeepSeek-R1-Distill-Qwen 7B and 14B; frontier models (DeepSeek V3.2, Claude Sonnet 4.5, Llama 3.3 70B) are tested behaviorally through assist-prefill. Model revisions, prompts, layer bands, and per-run configurations are in the appendix; code and data will be released.




Figure 1: The edit-and-continue protocol. A correct worked solution is truncated after one edited value line and the model continues. Three interventions target the edited value: changing the visible token, overwriting the hidden state at that position with the state a different value would have produced, and blocking attention from later positions to that position. The readout is whether the final answer follows the planted value.

Table 1: State patch on arithmetic chains at ten steps. Rates are conditioned on items where the text edit was followed (first column), so cross-model comparison of later columns should account for that rate. Brackets are 95 percent bootstrap intervals. The two reasoning-distilled models resolve the problem after their reasoning phase, which raises their random-control reversion; in a logged sample every reverting continuation re-derives from the problem’s starting value rather than correcting in place, and re-derivation can only produce the correct answer, so the planted-value column is unaffected.

model	text edit followed	restore correct	plant new value	random control	<i>n</i>
Qwen3-4B	0.98 [0.95, 1.00]	1.00 [1.00, 1.00]	1.00 [0.99, 1.00]	0.00 [0.00, 0.00]	106
Qwen2.5-7B-Instruct	0.85 [0.80, 0.90]	0.97 [0.92, 1.00]	0.74 [0.67, 0.81]	0.00 [0.00, 0.01]	93
Phi-3-medium	0.94 [0.90, 0.98]	0.98 [0.95, 1.00]	0.94 [0.91, 0.97]	0.00 [0.00, 0.00]	102
OLMo-2-7B	0.85 [0.79, 0.91]	0.88 [0.82, 0.94]	0.93 [0.88, 0.97]	0.00 [0.00, 0.00]	100
R1-distill-7B	0.60 [0.53, 0.66]	1.00 [0.99, 1.00]	0.76 [0.68, 0.84]	0.30 [0.22, 0.39]	65
R1-distill-14B	0.44 [0.36, 0.52]	0.99 [0.98, 1.00]	0.70 [0.61, 0.80]	0.21 [0.13, 0.29]	46

87 3 The written position holds a settable value register

88 Table 1 decomposes the effect of editing a written value. Restoring the correct hidden state while
89 the text still shows a corrupted value returns the answer to correct in 0.88 to 1.00 of cases. Planting
90 a third value that appears nowhere in the text redirects the answer to it in 0.70 to 1.00 of cases. The
91 size-matched random control sits at or below 0.01 for the four instruction-tuned models. Because
92 different planted values produce answers following those values, the value content is the causal
93 quantity, and because the planted value is a mid-chain intermediate the model never produced, the
94 patch does not inject the answer. The effect persists with task depth: on Qwen3-4B the planted value
95 is followed at 0.98, 0.97, and 0.95 for chains of 5, 20, and 40 steps. No fine-tuning, task-specific
96 training, or prompt engineering produces this behavior; it is how the pretrained models already work.

97 **The result is a property of written state, not of one template.** On the narrative task the pattern
98 replicates: a planted never-written count is followed at 0.83 [0.75, 0.90] on Qwen3-4B and 0.86
99 [0.79, 0.92] on Phi-3-medium, restoring the correct state returns the answer at 0.89 and 0.94, and
100 the random control is 0.00 in both ($n = 86$ and 89).

101 **Written positions are independently addressable.** In a task with two counters summed at the
102 end, planting a never-written value at one final position yields the corresponding mixed sum (0.94




Figure 2: Rate at which the answer follows the edited value under each blocked attention target. Blocking the value’s own position collapses following on every model; matched control blocks do not. Error bars are 95 percent bootstrap intervals.

Table 2: Attention block on arithmetic chains: rate at which the answer follows the edited value. OLMo-2-7B continues this format less reliably (low no-block rate) and shows the same contrast.

model	no block	value position	neighbor	random position	earlier value
Qwen3-4B	0.900	0.000	0.903	0.778	—
Qwen2.5-7B-Instruct	0.643	0.000	0.640	0.670	0.658
Phi-3-medium	0.880	0.000	0.863	0.865	0.850
OLMo-2-7B	0.333	0.005	0.297	0.342	0.305

and 0.93 on Qwen2.5-7B, 0.89 and 0.93 on Qwen3-4B). Planting both yields their joint sum at 0.90 and 0.88, at or above the product of the single-position rates (0.89 and 0.83). The random patch leaves the clean answer in place (0.96 and 0.95), and answers using only one of the two planted values occur at 0.01. The model reports a number assembled from two independently set registers.

4 Retrieval runs through attention to the written position

Blocking attention from all post-edit positions to the edited value’s position collapses the answer’s dependence on that value on every model tested (Table 2, Figure 2). Matched blocks of the adjacent words, of a random prompt position, or of an earlier step’s written value leave the dependence at its unblocked level. Blocking the operand the next step needs breaks the arithmetic itself, yielding answers that match neither the edit nor the correct value. The collapse holds across three families and on both task formats (natural-language task: 0.730 to 0.007 on Qwen3-4B, 0.720 to 0.003 on Phi-3-medium).

The block removes the value, not the ability to generate. Under the value-position block the unparseable-output rate stays at 0.000 (0.040 on OLMo-2-7B, equal to its unblocked rate). The model produces a fluent, parseable number that matches neither the edited value nor the correct one: it computes forward from a value it can no longer retrieve.

The retrieval happens before two-thirds depth. A layer-resolved version of the state patch locates the read in depth. On Qwen3-4B (36 layers, $n = 60$), overwriting the value’s hidden state at any single layer from 0 through 22 redirects the answer at 0.90 to 0.94, matching the full multi-layer patch (0.92), while single-layer patches at layer 24 and beyond redirect at 0.02 or less. The value can be rewritten anywhere upstream of its consumers and nowhere downstream of them, so retrieval completes by roughly two thirds of the network’s depth. A matching layer-resolved sweep of the attention block, and coarse sweeps on two further models, will appear in the final version; preliminary runs place the necessary attention layers in the middle third, consistent with the patch cliff.

127 **Reading is the preferred path, not the only one.** Replacing the value’s characters with unread-
 128 able dots, without touching attention, also removes following, and differently by model: the Qwen
 129 models reconstruct the value from the previous line and recover the correct answer in 0.60 to 0.73
 130 of cases, while Phi-3-medium does not recover (0.00).

131 5 What a calibrated J-space ablation does and does not disturb

132 Two results relate the written registers to the model’s internal capacity. First, the internal side cannot
 133 replace writing on these tasks: instructed to output only the final answer, with token counts confirm-
 134 ing nothing else was generated, models fail beyond one to two dependent steps (Qwen2.5-7B falls
 135 from 1.00 at $d = 1$ to 0.04 at $d = 4$; DeepSeek V3.2 at 671B reaches 0.23 at $d = 1$ and near zero
 136 above), and models write from the first step rather than upon overflow (0.93 to 1.00 of ground-truth
 137 intermediates appear in free traces at every difficulty, for every model from 1.5B to 671B).

138 Second, an ablation of J-space, the concept subspace a fitted Jacobian lens reads from the residual
 139 stream [Gurnee et al., 2026], leaves the externalized computation intact. At the strongest dose that
 140 preserves ordinary generation (calibrated and frozen on neutral text: next-token agreement 0.87,
 141 perplexity ratio 1.12), removing the top 16 active concepts at every position changes neither chain-
 142 of-thought accuracy nor bare-answer accuracy nor the amount written, on synthetic chains and on
 143 GSM8K, and does no more than a size-matched random-directions ablation. We read this as a
 144 boundary on the particular readout tested, for three reasons. The dose is capped at generation-
 145 preserving strength, the lens is one fitted artifact, and a written token’s contextual representation
 146 and its J-space component are not exclusive alternatives, since the lens samples the same residual
 147 states our patch overwrites. The J-space account itself reports that explicit chain of thought makes
 148 reasoning more robust to J-space ablation and interprets writing as externalizing state the model
 149 would otherwise carry internally [Gurnee et al., 2026]. Our result is consistent with that reading
 150 and sharpens it with a mechanism: on these tasks the externalized channel is where the serial state
 151 causally resides, and the tested lens readout does not capture it.

152 6 When the written value is read versus recomputed

153 Models read the register by default. On DeepSeek V3.2, an edited value the prompt merely implies
 154 is followed at 0.93, while an edited value the prompt generatively defines (the prompt states the two
 155 numbers whose sum becomes the running value) is reverted to correct at 0.43; with the defining
 156 sentence present but the edit two steps away, following returns to 0.96, so the reversion targets
 157 exactly the defined value. Claude Sonnet 4.5 recomputes more broadly, reverting most when a
 158 stated checkpoint is the value’s only source. Llama 3.3 70B follows edits at 0.97 whether or not the
 159 value is prompt-defined, despite solving the same chains from scratch, so the capacity to recompute
 160 does not imply recomputing. Qwen models at 4B and 7B never revert. These conditions come from
 161 separate experiments with different trace lengths and item sets; each within-experiment contrast is
 162 controlled, and we do not claim a single ordered scale across models or that reconstruction cost is
 163 the governing rule. Characterizing what governs the choice is open.

164 The same test bounds the mechanism on natural benchmarks: on GSM8K worked solutions, an
 165 edited intermediate is followed at 0.87 by Qwen3-4B and at 0.10 by V3.2, since a benchmark in-
 166 termediate is often one operation from numbers still in the prompt. Written values act as memory
 167 chiefly for state the model does not reconstruct on its own.

168 7 Related work

169 Our interventions build on causal tracing and activation patching [Meng et al., 2022] and causal
 170 abstraction [Geiger et al., 2021]; Heimersheim and Nanda [2024] caution that broad patches con-
 171 flate a component mattering with deleting information at a position, which the single-layer sweep
 172 in Section 4 addresses. That written intermediates can serve as memory is established: theoreti-
 173 cally [Merrill and Sabharwal, 2024], behaviorally [Lanchantin et al., 2023], and causally in specific
 174 settings [Zhu et al., 2025, Shih et al., 2026]. We add that the behavior is native to unmodified pre-
 175 trained models across four families and two formats, the attention-level retrieval route with matched
 176 controls, its depth localization, register composition, the qualified J-space boundary, and the read-

177 versus-recompute map. Faithfulness studies show stated reasoning and internal computation can
178 diverge [Turpin et al., 2023, Lanham et al., 2023]; filler-token results show computation can pass
179 through uninformative tokens [Pfau et al., 2024]; both bound our claims to state the model does not
180 regenerate internally.

181 8 Limitations

182 Tasks are numeric and count state tracking, so our claims concern serial task state rather than work-
183 ing memory in general; code, plans, and multi-entity state are untested. The J-space result bounds
184 one fitted lens at generation-preserving doses and does not measure the J-space content of the writ-
185 ten positions themselves. The attention-block layer sweep and the cross-model layer sweeps are
186 preliminary. White-box evidence is at 4B to 14B; frontier evidence is behavioral. The read-versus-
187 recompute results support within-experiment contrasts, and the only patch control is size-matched
188 noise; a structured control planting a valid activation for a different quantity would further isolate
189 value content.

190 9 Conclusion

191 On serial state-tracking tasks, unmodified pretrained language models spontaneously use the hidden
192 states at written intermediate-value positions as compositional registers: settable by intervention, re-
193 trieved through attention before two-thirds depth, independently addressable, and decisive for the an-
194 swer. A calibrated J-space ablation leaves this externalized computation undisturbed, which bounds
195 that readout on these tasks. For oversight the implication cuts both ways. The written positions are
196 a causally used memory channel, so the trace exposes the symbols the computation operates on; and
197 identical visible text can coexist with different stored values and different downstream behavior, so
198 the text does not guarantee the hidden content at those positions. Trace monitoring observes the
199 channel, and verifying the state behind it is a separate problem.

200 References

- 201 Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. Causal abstractions of
202 neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
203 arXiv:2106.02997.
- 204 Wes Gurnee et al. Reading concepts from the residual stream with jacobian lenses. *arXiv preprint*,
205 2026.
- 206 Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv preprint*
207 *arXiv:2404.15255*, 2024.
- 208 Jack Lanchantin, Shubham Toshniwal, Jason Weston, Arthur Szlam, and Sainbayar Sukhbaatar.
209 Learning to reason and memorize with self-notes. In *Advances in Neural Information Processing*
210 *Systems (NeurIPS)*, 2023. arXiv:2305.00833.
- 211 Tamera Lanham et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint*
212 *arXiv:2307.13702*, 2023.
- 213 Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual
214 associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
215 arXiv:2202.05262.
- 216 William Merrill and Ashish Sabharwal. The expressive power of transformers with chain of thought.
217 In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.07923.
- 218 Jacob Pfau, William Merrill, and Samuel R. Bowman. Let’s think dot by dot: Hidden compu-
219 tation in transformer language models. In *Conference on Language Modeling (COLM)*, 2024.
220 arXiv:2404.15758.
- 221 Andy Shih et al. Do models read what they write? causal registers in scratchpad reasoning. *arXiv*
222 *preprint arXiv:2606.29522*, 2026.

- 223 Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don't always
224 say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in*
225 *Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.04388.
- 226 Fangwei Zhu et al. Chain-of-thought tokens are computer program variables. *arXiv preprint*
227 *arXiv:2505.04955*, 2025.