

- 228 Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don't always
229 say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in*
230 *Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.04388.
- 231 Fangwei Zhu et al. Chain-of-thought tokens are computer program variables. *arXiv preprint*
232 *arXiv:2505.04955*, 2025.