
The Chain of Thought as Compositional Memory in Unmodified Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Unmodified pretrained language models use the hidden states at written
2 intermediate-value tokens as compositional registers for serial task state. Over-
3 writing the hidden state at one written value’s position with the state a different
4 number would have produced, while leaving the visible text unchanged, redirects
5 the final answer to the never-written number in 0.70 to 1.00 of cases across six
6 models from four families on arithmetic and narrative state-tracking tasks; a size-
7 matched random perturbation redirects at most 0.01 for instruction-tuned models.
8 The value is retrieved through attention to its own position: blocking that attention
9 collapses the answer’s dependence on the value to at most 0.007 on four models,
10 matched blocks of neighboring, random, and earlier-value positions change nothing,
11 and the model continues to produce fluent parseable answers. Layer-resolved
12 sweeps time the retrieval: planting through a single layer succeeds at any depth
13 up to two thirds of the network and fails after, and blocking the attention only at
14 the middle third of layers reproduces the full collapse, so the value is read by at-
15 tention in the middle third and is causally inert afterward. Two positions function
16 as independent stores: values planted at both compose into their sum, a number
17 appearing nowhere in the input, at 0.88 to 0.90, and answers drawing on only one
18 planted value occur at 0.01. A calibrated ablation of the concepts a fitted Jacobian
19 lens reads from the residual stream (J-space), at the strongest dose preserving ordi-
20 nary generation, leaves this externalized computation intact, bounding what that
21 readout captures rather than locating serial state outside the workspace it samples.
22 These results carry a caution for oversight: the written positions are a causally
23 used memory channel, and identical visible text can coexist with different stored
24 values and different downstream behavior, so a trace exposes the symbols of the
25 computation without guaranteeing their hidden content.

26 1 Introduction

27 When a pretrained language model tracks serial state through a chain of thought, the hidden state at
28 each written value’s token is a memory register: settable by intervention, retrieved through attention
29 to that position, and independently addressable, with the stored value rather than the visible text
30 determining the answer. We establish this with causal interventions on six unmodified models from
31 four families (Qwen, Phi, OLMo, DeepSeek-R1-Distill), on two task families with exact ground
32 truth for every intermediate value, and we locate the retrieval in the network’s depth.

33 Prior work anticipated pieces of this picture without establishing the mechanism in the wild. Ex-
34 pressivity theory treats generated tokens as the recurrent state of an otherwise fixed-depth compu-
35 tation [Merrill and Sabharwal, 2024], Self-Notes demonstrates writing as working memory behav-
36 iorally [Lanchantin et al., 2023], Zhu et al. [2025] show that intermediate-result tokens store muta-




Figure 1: The edit-and-continue protocol. A correct worked solution is truncated after one edited value line and the model continues. Three interventions target the edited value: changing the visible token, overwriting the hidden state at that position with the state a different value would have produced, and blocking attention from later positions to that position. The readout is whether the final answer follows the planted value.

ble values on multiplication and dynamic-programming tasks, and Shih et al. [2026] patch written scratchpad representations on a fine-tuned model. None of these show that unmodified pretrained models do this natively, identify the retrieval route, test whether written stores compose, or map when a model reads a written value rather than recomputing it. These questions are this paper’s results:

1. **The hidden state at a written value’s position is a value register, natively.** Planting a never-written value into the residual stream at that position, text unchanged, redirects the answer to the planted value at 0.70 to 1.00 across all six models; restoring the correct state under corrupted text reverts the answer at 0.88 to 1.00; size-matched random noise redirects at most 0.01 on instruction-tuned models (Section 3).
2. **Retrieval runs through attention to the written position, before two-thirds depth.** Blocking attention from later positions to that one position collapses the answer’s dependence on the value to at most 0.007 on four models in three families, with matched control blocks at the unblocked level and fluency unchanged. Planting through a single layer succeeds at any layer up to two thirds of the network and fails at deeper layers, and blocking attention only at the middle third of layers reproduces the full collapse, so the value is consumed by attention in the middle third (Section 4).
3. **Written positions are independently addressable.** Two never-written values planted at two positions yield their sum, absent from the input, at 0.88 to 0.90; answers using only one of the two occur at 0.01 (Section 3).
4. **A calibrated J-space ablation does not disturb the externalized computation.** Removing the top concepts a fitted Jacobian lens reads from the residual stream, at the strongest dose that preserves ordinary generation, changes neither accuracy nor what gets written. This bounds the particular readout tested; it does not place serial state outside the workspace the lens samples (Section 5).
5. **Reading is the default; recomputation is model-specific.** Models follow an edited written value even when the prompt implies the correct one. DeepSeek V3.2 recomputes only values the prompt generatively defines, and precisely those; Claude Sonnet 4.5 also uses stated checkpoints; Llama 3.3 70B recomputes almost nothing despite solving the same chains from scratch (Section 6).

2 Tasks and protocol

Both task families give exact ground truth for every intermediate value, so the downstream effect of any edit is computable. Arithmetic chains start from a three-digit number and apply d signed two-digit additions and subtractions. Narrative state-tracking problems describe a character whose count

Table 1: State patch on arithmetic chains at ten steps. Rates are conditioned on items where the text edit was followed (first column), so cross-model comparison of later columns should account for that rate. Brackets are 95 percent bootstrap intervals. The two reasoning-distilled models resolve the problem after their reasoning phase, which raises their random-control reversion; in a logged sample every reverting continuation re-derives from the problem’s starting value rather than correcting in place, and re-derivation can only produce the correct answer, so the planted-value column is unaffected.

model	text edit followed	restore correct	plant new value	random control	n
Qwen3-4B	0.98 [0.95, 1.00]	1.00 [1.00, 1.00]	1.00 [0.99, 1.00]	0.00 [0.00, 0.00]	106
Qwen2.5-7B-Instruct	0.85 [0.80, 0.90]	0.97 [0.92, 1.00]	0.74 [0.67, 0.81]	0.00 [0.00, 0.01]	93
Phi-3-medium	0.94 [0.90, 0.98]	0.98 [0.95, 1.00]	0.94 [0.91, 0.97]	0.00 [0.00, 0.00]	102
OLMo-2-7B	0.85 [0.79, 0.91]	0.88 [0.82, 0.94]	0.93 [0.88, 0.97]	0.00 [0.00, 0.00]	100
R1-distill-7B	0.60 [0.53, 0.66]	1.00 [0.99, 1.00]	0.76 [0.68, 0.84]	0.30 [0.22, 0.39]	65
R1-distill-14B	0.44 [0.36, 0.52]	0.99 [0.98, 1.00]	0.70 [0.61, 0.80]	0.21 [0.13, 0.29]	46

of an item changes through varied events, with varied names, items, verbs, and sentence forms. In both, the model writes one line per step giving the running value.

The protocol takes a correct worked solution, changes the value at one step j , truncates immediately after the edited line, and lets the model continue (Figure 1). A no-edit truncation measures the sampling noise floor, 0.00 throughout. Generation uses temperature 0.6 and top- p 0.95 with five samples per item; unparseable outputs are counted separately. Intervals are 95 percent bootstrap intervals over items; the effects that carry the argument exceed their controls by far larger margins. White-box experiments use Qwen3-4B, Qwen2.5-7B-Instruct, Phi-3-medium, OLMo-2-7B-Instruct, and DeepSeek-R1-Distill-Qwen 7B and 14B; frontier models (DeepSeek V3.2, Claude Sonnet 4.5, Llama 3.3 70B) are tested behaviorally through assistant prefill. Model revisions, prompts, layer bands, and configurations are in the appendix; code and data will be released.

3 The written position holds a settable value register

Table 1 decomposes the effect of editing a written value. Restoring the correct hidden state while the text still shows a corrupted value returns the answer to correct in 0.88 to 1.00 of cases. Planting a third value that appears nowhere in the text redirects the answer to it in 0.70 to 1.00 of cases. The size-matched random control sits at or below 0.01 for the four instruction-tuned models. Because different planted values produce answers following those values, the value content is the causal quantity, and because the planted value is a mid-chain intermediate the model never produced, the patch does not inject the answer. The effect persists with task depth: on Qwen3-4B the planted value is followed at 0.98, 0.97, and 0.95 for chains of 5, 20, and 40 steps. No fine-tuning, task-specific training, or prompt engineering produces this behavior; it is how the pretrained models already work.

The result is a property of written state, not of one template. On the narrative task the pattern replicates: a planted never-written count is followed at 0.83 [0.75, 0.90] on Qwen3-4B and 0.86 [0.79, 0.92] on Phi-3-medium, restoring the correct state returns the answer at 0.89 and 0.94, and the random control is 0.00 in both ($n = 86$ and 89).

Written positions are independently addressable. In a task with two counters summed at the end, planting a never-written value at one final position yields the corresponding mixed sum (0.94 and 0.93 on Qwen2.5-7B, 0.89 and 0.93 on Qwen3-4B). Planting both yields their joint sum at 0.90 and 0.88, at or above the product of the single-position rates (0.89 and 0.83). The random patch leaves the clean answer in place (0.96 and 0.95), and answers using only one of the two planted values occur at 0.01. The model reports a number assembled from two independently set registers.

4 Retrieval runs through attention to the written position

Blocking attention from all post-edit positions to the edited value’s position collapses the answer’s dependence on that value on every model tested (Table 2, Figure 2). Matched blocks of the adjacent




Figure 2: Rate at which the answer follows the edited value under each blocked attention target. Blocking the value’s own position collapses following on every model; matched control blocks do not. Error bars are 95 percent bootstrap intervals.

Table 2: Attention block on arithmetic chains: rate at which the answer follows the edited value. OLMo-2-7B continues this format less reliably (low no-block rate) and shows the same contrast.

model	no block	value position	neighbor	random position	earlier value
Qwen3-4B	0.900	0.000	0.903	0.778	—
Qwen2.5-7B-Instruct	0.643	0.000	0.640	0.670	0.658
Phi-3-medium	0.880	0.000	0.863	0.865	0.850
OLMo-2-7B	0.333	0.005	0.297	0.342	0.305

words, of a random prompt position, or of an earlier step’s written value leave the dependence at its unblocked level. Blocking the operand the next step needs breaks the arithmetic itself, yielding answers that match neither the edit nor the correct value. The collapse holds across three families and on both task formats (natural-language task: 0.730 to 0.007 on Qwen3-4B, 0.720 to 0.003 on Phi-3-medium).

The block removes the value, not the ability to generate. Under the value-position block the unparseable-output rate stays at 0.000 (0.040 on OLMo-2-7B, equal to its unblocked rate). The model produces a fluent, parseable number that matches neither the edited value nor the correct one: it computes forward from a value it can no longer retrieve.

Retrieval occurs in the middle third of the network. Layer-resolved versions of both interventions locate the read in depth on Qwen3-4B (36 layers). Overwriting the value’s hidden state at any single layer from 0 through 22 redirects the answer at 0.90 to 0.94, matching the full multi-layer patch (0.92), while single-layer patches at layer 24 and beyond redirect at 0.02 or less ($n = 60$): the register can be rewritten anywhere upstream of its consumers and nowhere downstream of them. Blocking the value-position attention only at the middle third of layers (12 through 23) collapses following from 0.857 to 0.025, blocking only the early or late third leaves it unchanged (0.857 and 0.863), and no single layer’s block matters (0.84 to 0.89 at every layer, $n = 80$): the read is distributed redundantly across the middle third’s layers and confined to it. The two sweeps agree on the same boundary, near layer 23: the value is written into the position’s residual stream from early depth, retrieved by attention across the middle third, and causally inert afterward. Coarse sweeps on two further models appear in the appendix.

Reading is the preferred path, not the only one. Replacing the value’s characters with unreadable dots, without touching attention, also removes following, and differently by model: the Qwen models reconstruct the value from the previous line and recover the correct answer in 0.60 to 0.73 of cases, while Phi-3-medium does not recover (0.00).

5 What a calibrated J-space ablation does and does not disturb

Two results relate the written registers to the model’s internal capacity. First, the internal side cannot replace writing on these tasks: instructed to output only the final answer, with token counts confirming nothing else was generated, models fail beyond one to two dependent steps (Qwen2.5-7B falls from 1.00 at $d = 1$ to 0.04 at $d = 4$; DeepSeek V3.2 at 671B reaches 0.23 at $d = 1$ and near zero above), and models write from the first step rather than upon overflow (0.93 to 1.00 of ground-truth intermediates appear in free traces at every difficulty, for every model from 1.5B to 671B).

Second, an ablation of J-space, the concept subspace a fitted Jacobian lens reads from the residual stream [Gurnee et al., 2026], leaves the externalized computation intact. At the strongest dose that preserves ordinary generation (calibrated and frozen on neutral text: next-token agreement 0.87, perplexity ratio 1.12), removing the top 16 active concepts at every position changes neither chain-of-thought accuracy nor bare-answer accuracy nor the amount written, on synthetic chains and on GSM8K, and does no more than a size-matched random-directions ablation. We read this as a boundary on the particular readout tested, for three reasons. The dose is capped at generation-preserving strength, the lens is one fitted artifact, and a written token’s contextual representation and its J-space component are not exclusive alternatives, since the lens samples the same residual states our patch overwrites. The J-space account itself reports that explicit chain of thought makes reasoning more robust to J-space ablation and interprets writing as externalizing state the model would otherwise carry internally [Gurnee et al., 2026]. Our result is consistent with that reading and sharpens it with a mechanism: on these tasks the externalized channel is where the serial state causally resides, and the tested lens readout does not capture it.

6 When the written value is read versus recomputed

Models read the register by default. On DeepSeek V3.2, an edited value the prompt merely implies is followed at 0.93, while an edited value the prompt generatively defines (the prompt states the two numbers whose sum becomes the running value) is reverted to correct at 0.43; with the defining sentence present but the edit two steps away, following returns to 0.96, so the reversion targets exactly the defined value. Claude Sonnet 4.5 recomputes more broadly, reverting most when a stated checkpoint is the value’s only source. Llama 3.3 70B follows edits at 0.97 whether or not the value is prompt-defined, despite solving the same chains from scratch, so the capacity to recompute does not imply recomputing. Qwen models at 4B and 7B never revert. These conditions come from separate experiments with different trace lengths and item sets; each within-experiment contrast is controlled, and we do not claim a single ordered scale across models or that reconstruction cost is the governing rule. Characterizing what governs the choice is open.

The same test bounds the mechanism on natural benchmarks: on GSM8K worked solutions, an edited intermediate is followed at 0.87 by Qwen3-4B and at 0.10 by V3.2, since a benchmark intermediate is often one operation from numbers still in the prompt. Written values act as memory chiefly for state the model does not reconstruct on its own.

7 Related work

Our interventions build on causal tracing and activation patching [Meng et al., 2022] and causal abstraction [Geiger et al., 2021]; Heimersheim and Nanda [2024] caution that broad patches conflate a component mattering with deleting information at a position, which the single-layer sweep in Section 4 addresses. That written intermediates can serve as memory is established: theoretically [Merrill and Sabharwal, 2024], behaviorally [Lanchantin et al., 2023], and causally in specific settings [Zhu et al., 2025, Shih et al., 2026]. We add that the behavior is native to unmodified pre-trained models across four families and two formats, the attention-level retrieval route with matched controls, its depth localization, register composition, the qualified J-space boundary, and the read-versus-recompute map. Faithfulness studies show stated reasoning and internal computation can diverge [Turpin et al., 2023, Lanham et al., 2023]; filler-token results show computation can pass through uninformative tokens [Pfau et al., 2024]; both bound our claims to state the model does not regenerate internally.

8 Limitations

Tasks are numeric and count state tracking, so our claims concern serial task state rather than working memory in general; code, plans, and multi-entity state are untested. The J-space result bounds one fitted lens at generation-preserving doses and does not measure the J-space content of the written positions themselves. The layer sweeps are established on one model; the cross-model sweeps are coarse. White-box evidence is at 4B to 14B; frontier evidence is behavioral. The read-versus-recompute results support within-experiment contrasts, and the only patch control is size-matched noise; a structured control planting a valid activation for a different quantity would further isolate value content.

9 Conclusion

On serial state-tracking tasks, unmodified pretrained language models spontaneously use the hidden states at written intermediate-value positions as compositional registers: settable by intervention, retrieved through attention before two-thirds depth, independently addressable, and decisive for the answer. A calibrated J-space ablation leaves this externalized computation undisturbed, which bounds that readout on these tasks. For oversight the implication cuts both ways. The written positions are a causally used memory channel, so the trace exposes the symbols the computation operates on; and identical visible text can coexist with different stored values and different downstream behavior, so the text does not guarantee the hidden content at those positions. Trace monitoring observes the channel, and verifying the state behind it is a separate problem.

References

- Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. Causal abstractions of neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021. arXiv:2106.02997.
- Wes Gurnee et al. Reading concepts from the residual stream with jacobian lenses. *arXiv preprint*, 2026.
- Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv preprint arXiv:2404.15255*, 2024.
- Jack Lanchantin, Shubham Toshniwal, Jason Weston, Arthur Szlam, and Sainbayar Sukhbaatar. Learning to reason and memorize with self-notes. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.00833.
- Tamera Lanham et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2202.05262.
- William Merrill and Ashish Sabharwal. The expressive power of transformers with chain of thought. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.07923.
- Jacob Pfau, William Merrill, and Samuel R. Bowman. Let’s think dot by dot: Hidden computation in transformer language models. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2404.15758.
- Andy Shih et al. Do models read what they write? causal registers in scratchpad reasoning. *arXiv preprint arXiv:2606.29522*, 2026.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.04388.
- Fangwei Zhu et al. Chain-of-thought tokens are computer program variables. *arXiv preprint arXiv:2505.04955*, 2025.