
No Workspace to Spill From: The Chain of Thought Is the Working Memory of Serial Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Language models cannot hold serial reasoning state internally, and the chain of
2 thought is where that state actually lives. A common reading of J-space, the con-
3 cept workspace a fitted Jacobian lens reads from the residual stream, is that rea-
4 soning is carried internally and writing relieves the workspace by externalizing
5 state it would otherwise hold. That account implies an internal alternative to writ-
6 ing, and three measurements find none. Instructed to answer chained arithmetic
7 without writing, models from 1.5B to 671B parameters collapse beyond one to
8 two dependent steps, so a 671B model cannot chain two additions in its head.
9 Models write every intermediate value from the first step at every difficulty, so
10 externalization begins at no capacity threshold. The allocation between the two
11 sides never moves: damaging the workspace does not increase writing, capping
12 writing does not move computation inward, and ablating the model’s active J-
13 space concepts at the strongest dose that preserves ordinary generation changes
14 nothing, even at depths where the workspace does perform the arithmetic. We
15 then show how this only memory works, with causal interventions on six unmod-
16 ified models from four families on arithmetic and narrative state-tracking tasks.
17 Each written value is held at its own token position as a register: overwriting the
18 hidden state there with the state a different number would have produced, text un-
19 changed, redirects the final answer to that never-written number in 0.70 to 1.00
20 of cases, while size-matched random noise redirects at most 0.01. The register is
21 read through attention to its position (blocking it collapses dependence to at most
22 0.007 with matched controls unchanged), the read completes by two thirds of the
23 network’s depth, and two positions act as independent stores whose planted values
24 compose into a sum appearing nowhere in the input (0.88 to 0.90). For oversight,
25 identical visible text can carry different stored values with different downstream
26 behavior, so a chain of thought exposes the symbols a computation operates on
27 without guaranteeing the hidden content behind them.

28 1 Introduction

29 The chain of thought is the working memory of serial reasoning in language models, and there is no
30 internal workspace standing behind it. We establish both halves of this claim with causal interven-
31 tions on six unmodified models (used as released, with no fine-tuning or task-specific training) from
32 four families (Qwen, Phi, OLMo, DeepSeek-R1-Distill), on two task families with exact ground
33 truth for every intermediate value.

34 The first half contradicts a natural reading of J-space, the concept workspace a fitted Jacobian lens
35 reads from the residual stream [Gurnee et al., 2026]: that reasoning state is carried internally and




Figure 1: The edit-and-continue protocol. A correct worked solution is truncated after one edited value line and the model continues. Three interventions target the edited value: changing the visible token, overwriting the hidden state at that position with the state a different value would have produced, and blocking attention from later positions to that position. The readout is whether the final answer follows the planted value.

36 writing externalizes state the workspace would otherwise hold. An internal carrier of serial state
 37 should show up in three places, and it shows up in none of them.

- 38 1. **No fallback.** Instructed to answer chained arithmetic with no writing (token counts confirm
 39 the channel is closed), models collapse beyond one to two dependent steps; DeepSeek V3.2,
 40 at 671B parameters, answers 0.23 at one step and near zero above (Section 3).
- 41 2. **No threshold.** Models write 0.93 to 1.00 of ground-truth intermediates from one-step
 42 problems onward, at every scale from 1.5B to 671B (Section 3).
- 43 3. **No reallocation.** Damaging the workspace does not increase writing, capping writing does
 44 not move computation inward, and ablating the top active J-space concepts at a generation-
 45 preserving dose changes neither accuracy nor what gets written (Section 3).

46 The second half is the mechanism of the memory that remains.

- 47 4. **Registers.** The hidden state at a written value’s position is a settable value register: planting
 48 a never-written value there, text unchanged, redirects the answer at 0.70 to 1.00 across all
 49 six models; restoring the correct state under corrupted text reverts it at 0.88 to 1.00; size-
 50 matched random noise redirects at most 0.01 on instruction-tuned models (Section 4).
- 51 5. **Retrieval.** The register is read through attention to its position: blocking that attention
 52 collapses dependence to at most 0.007 on four models in three families, with neighbor,
 53 random, and earlier-value blocks at the unblocked level and fluency unchanged. Layer
 54 sweeps time the read: single-layer patches succeed at any depth up to two thirds of the
 55 network and fail after, and blocking attention only at the middle third reproduces the full
 56 collapse (Section 5).
- 57 6. **Composition, and when the register is trusted.** Two positions are independently ad-
 58 dressable, composing planted values into their sum (0.88 to 0.90). Models read the register
 59 by default; which prompt evidence triggers recomputation instead is model-specific (Sec-
 60 tions 4 and 6).

61 Prior work anticipated the memory role of written tokens: theoretically [Merrill and Sabharwal,
 62 2024], behaviorally [Lanchantin et al., 2023], and causally in specific settings, on a fine-tuned
 63 model [Shih et al., 2026] and on multiplication and dynamic-programming tasks [Zhu et al., 2025].
 64 None of it tests whether an internal alternative exists, shows the behavior is native to unmodified
 65 pretrained models, identifies the retrieval route, or tests composition. Those are this paper’s results.

66 2 Tasks and protocol

67 Both task families give exact ground truth for every intermediate value, so the downstream effect of
68 any edit is computable. Arithmetic chains start from a three-digit number and apply d signed two-
69 digit additions and subtractions. Narrative state-tracking problems describe a character whose count
70 of an item changes through varied events, with varied names, items, verbs, and sentence forms. In
71 both, the model writes one line per step giving the running value.

72 The protocol takes a correct worked solution, changes the value at one step j , truncates immedi-
73 ately after the edited line, and lets the model continue (Figure 1). A no-edit truncation measures
74 the sampling noise floor, 0.00 throughout. Generation uses temperature 0.6 and top- p 0.95 with five
75 samples per item; unparseable outputs are counted separately. Intervals are 95 percent bootstrap
76 intervals over items; the effects that carry the argument exceed their controls by far larger mar-
77 gins. White-box experiments use Qwen3-4B, Qwen2.5-7B-Instruct, Phi-3-medium, OLMo-2-7B-
78 Instruct, and DeepSeek-R1-Distill-Qwen 7B and 14B; frontier models (DeepSeek V3.2, Claude Son-
79 net 4.5, Llama 3.3 70B) are tested behaviorally through assistant prefill. Model revisions, prompts,
80 layer bands, and configurations are in the appendix; code and data will be released.

81 3 Serial state has no internal home

82 Three measurements look for an internal carrier of chained reasoning state and find none. First,
83 capacity: instructed to output only the final answer, with token counts confirming nothing else was
84 generated, models fail beyond one to two dependent steps (Qwen2.5-7B: 1.00 at $d = 1$, 0.04 at
85 $d = 4$; DeepSeek V3.2 at 671B: 0.23 at $d = 1$, near zero above), and the values used are random,
86 so nothing can be recalled rather than computed. Second, the write policy: models write 0.93 to
87 1.00 of ground-truth intermediates in free generation from $d = 1$ through $d = 64$, at every scale
88 from 1.5B to 671B, so writing does not begin at any capacity threshold. Third, reallocation: a
89 residual-stream lesion strong enough to cross the accuracy cliff makes traces disintegrate rather than
90 expand; hard token budgets compress wording up to 2.5 times while keeping 0.98 to 1.00 of values,
91 and below roughly twelve tokens per step the model truncates and fails rather than holding values
92 internally; and ablating J-space, the concept subspace a fitted Jacobian lens reads from the residual
93 stream [Gurnee et al., 2026], changes nothing at the strongest dose that preserves ordinary generation
94 (calibrated and frozen on neutral text: next-token agreement 0.87, perplexity ratio 1.12). The J-space
95 ablation removes the top 16 active concepts at every position and leaves chain-of-thought accuracy,
96 bare-answer accuracy, and the amount written unchanged, on synthetic chains and on GSM8K, doing
97 no more than a size-matched random-directions ablation. The null holds at one to two steps, where
98 models do perform the arithmetic internally, so the load-bearing content there lies outside the top
99 active concepts at generation-preserving strength.

100 This bounds the particular readout tested rather than the workspace as a whole: the dose is capped
101 at generation-preserving strength, the lens is one fitted artifact, and a written token’s contextual
102 representation and its J-space component are not exclusive alternatives, since the lens samples the
103 same residual states our patches overwrite. The J-space account itself reports that explicit chain of
104 thought makes reasoning more robust to J-space ablation [Gurnee et al., 2026]. Our measurements
105 sharpen that observation into a stronger statement: writing is not relieving an internal store, because
106 closing the written channel reveals no internal store to relieve, at any scale tested.

107 4 The written position holds a settable value register

108 Table 1 decomposes the effect of editing a written value. Restoring the correct hidden state while
109 the text still shows a corrupted value returns the answer to correct in 0.88 to 1.00 of cases. Planting
110 a third value that appears nowhere in the text redirects the answer to it in 0.70 to 1.00 of cases. The
111 size-matched random control sits at or below 0.01 for the four instruction-tuned models. Because
112 different planted values produce answers following those values, the value content is the causal
113 quantity, and because the planted value is a mid-chain intermediate the model never produced, the
114 patch does not inject the answer. The effect persists with task depth: on Qwen3-4B the planted value
115 is followed at 0.98, 0.97, and 0.95 for chains of 5, 20, and 40 steps. No fine-tuning, task-specific
116 training, or prompt engineering produces this behavior; it is how the pretrained models already work.

Table 1: State patch on arithmetic chains at ten steps. Rates are conditioned on items where the text edit was followed (first column), so cross-model comparison of later columns should account for that rate. Brackets are 95 percent bootstrap intervals. The two reasoning-distilled models resolve the problem after their reasoning phase, which raises their random-control reversion; in a logged sample every reverting continuation re-derives from the problem’s starting value rather than correcting in place, and re-derivation can only produce the correct answer, so the planted-value column is unaffected.

model	text edit followed	restore correct	plant new value	random control	n
Qwen3-4B	0.98 [0.95, 1.00]	1.00 [1.00, 1.00]	1.00 [0.99, 1.00]	0.00 [0.00, 0.00]	106
Qwen2.5-7B-Instruct	0.85 [0.80, 0.90]	0.97 [0.92, 1.00]	0.74 [0.67, 0.81]	0.00 [0.00, 0.01]	93
Phi-3-medium	0.94 [0.90, 0.98]	0.98 [0.95, 1.00]	0.94 [0.91, 0.97]	0.00 [0.00, 0.00]	102
OLMo-2-7B	0.85 [0.79, 0.91]	0.88 [0.82, 0.94]	0.93 [0.88, 0.97]	0.00 [0.00, 0.00]	100
R1-distill-7B	0.60 [0.53, 0.66]	1.00 [0.99, 1.00]	0.76 [0.68, 0.84]	0.30 [0.22, 0.39]	65
R1-distill-14B	0.44 [0.36, 0.52]	0.99 [0.98, 1.00]	0.70 [0.61, 0.80]	0.21 [0.13, 0.29]	46




Figure 2: Rate at which the answer follows the edited value under each blocked attention target. Blocking the value’s own position collapses following on every model; matched control blocks do not. Error bars are 95 percent bootstrap intervals.

117 **The result is a property of written state, not of one template.** On the narrative task the pattern
118 replicates: a planted never-written count is followed at 0.83 [0.75, 0.90] on Qwen3-4B and 0.86
119 [0.79, 0.92] on Phi-3-medium, restoring the correct state returns the answer at 0.89 and 0.94, and
120 the random control is 0.00 in both ($n = 86$ and 89).

121 **Written positions are independently addressable.** In a task with two counters summed at the
122 end, planting a never-written value at one final position yields the corresponding mixed sum (0.94
123 and 0.93 on Qwen2.5-7B, 0.89 and 0.93 on Qwen3-4B). Planting both yields their joint sum at 0.90
124 and 0.88, at or above the product of the single-position rates (0.89 and 0.83). The random patch
125 leaves the clean answer in place (0.96 and 0.95), and answers using only one of the two planted
126 values occur at 0.01. The model reports a number assembled from two independently set registers.

127 5 Retrieval runs through attention to the written position

128 Blocking attention from all post-edit positions to the edited value’s position collapses the answer’s
129 dependence on that value on every model tested (Table 2, Figure 2). Matched blocks of the adjacent
130 words, of a random prompt position, or of an earlier step’s written value leave the dependence at
131 its unblocked level. Blocking the operand the next step needs breaks the arithmetic itself, yielding
132 answers that match neither the edit nor the correct value. The collapse holds across three families
133 and on both task formats (natural-language task: 0.730 to 0.007 on Qwen3-4B, 0.720 to 0.003 on
134 Phi-3-medium).

135 **The block removes the value, not the ability to generate.** Under the value-position block the
136 unparseable-output rate stays at 0.000 (0.040 on OLMo-2-7B, equal to its unblocked rate). The

Table 2: Attention block on arithmetic chains: rate at which the answer follows the edited value. OLMo-2-7B continues this format less reliably (low no-block rate) and shows the same contrast.

model	no block	value position	neighbor	random position	earlier value
Qwen3-4B	0.900	0.000	0.903	0.778	—
Qwen2.5-7B-Instruct	0.643	0.000	0.640	0.670	0.658
Phi-3-medium	0.880	0.000	0.863	0.865	0.850
OLMo-2-7B	0.333	0.005	0.297	0.342	0.305

model produces a fluent, parseable number that matches neither the edited value nor the correct one: it computes forward from a value it can no longer retrieve.

Retrieval occurs in the middle third of the network. Layer-resolved versions of both interventions locate the read in depth on Qwen3-4B (36 layers). Overwriting the value’s hidden state at any single layer from 0 through 22 redirects the answer at 0.90 to 0.94, matching the full multi-layer patch (0.92), while single-layer patches at layer 24 and beyond redirect at 0.02 or less ($n = 60$): the register can be rewritten anywhere upstream of its consumers and nowhere downstream of them. Blocking the value-position attention only at the middle third of layers (12 through 23) collapses following from 0.857 to 0.025, blocking only the early or late third leaves it unchanged (0.857 and 0.863), and no single layer’s block matters (0.84 to 0.89 at every layer, $n = 80$): the read is distributed redundantly across the middle third’s layers and confined to it. The two sweeps agree on the same boundary, near layer 23: the value is written into the position’s residual stream from early depth, retrieved by attention across the middle third, and causally inert afterward. Coarse sweeps on two further models (appendix) show the structure generalizes while the boundary’s depth is model-specific: on Phi-3-medium, early- and mid-third patches redirect at 0.89 and 0.87 and late-third patches at 0.007, matching the Qwen3-4B shape, while on Qwen2.5-7B patches in every third redirect at 0.70 to 0.73, so its reads extend into the final third of its 28 layers.

Reading is the preferred path, not the only one. Replacing the value’s characters with unreadable dots, without touching attention, also removes following, and differently by model: the Qwen models reconstruct the value from the previous line and recover the correct answer in 0.60 to 0.73 of cases, while Phi-3-medium does not recover (0.00).

6 When the written value is read versus recomputed

Models read the register by default. On DeepSeek V3.2, an edited value the prompt merely implies is followed at 0.93, while an edited value the prompt generatively defines (the prompt states the two numbers whose sum becomes the running value) is reverted to correct at 0.43; with the defining sentence present but the edit two steps away, following returns to 0.96, so the reversion targets exactly the defined value. Claude Sonnet 4.5 recomputes more broadly, reverting most when a stated checkpoint is the value’s only source. Llama 3.3 70B follows edits at 0.97 whether or not the value is prompt-defined, despite solving the same chains from scratch, so the capacity to recompute does not imply recomputing. Qwen models at 4B and 7B never revert. These conditions come from separate experiments with different trace lengths and item sets; each within-experiment contrast is controlled, and we do not claim a single ordered scale across models or that reconstruction cost is the governing rule. Characterizing what governs the choice is open.

The same test bounds the mechanism on natural benchmarks: on GSM8K worked solutions, an edited intermediate is followed at 0.87 by Qwen3-4B and at 0.10 by V3.2, since a benchmark intermediate is often one operation from numbers still in the prompt. Written values act as memory chiefly for state the model does not reconstruct on its own.

7 Related work

Our interventions build on causal tracing and activation patching [Meng et al., 2022] and causal abstraction [Geiger et al., 2021]; Heimersheim and Nanda [2024] caution that broad patches conflate a component mattering with deleting information at a position, which the single-layer sweep

in Section 5 addresses. The memory role of written tokens in prior work is discussed in the introduction [Merrill and Sabharwal, 2024, Lanchantin et al., 2023, Zhu et al., 2025, Shih et al., 2026]. Faithfulness studies [Turpin et al., 2023, Lanham et al., 2023] and filler-token results [Pfau et al., 2024] bound our claims to state the model does not regenerate internally.

8 Limitations

Tasks are numeric and count state tracking, so our claims concern serial task state rather than working memory in general; code, plans, and multi-entity state are untested. The J-space result bounds one fitted lens at generation-preserving doses and does not measure the J-space content of the written positions themselves. The fine layer sweeps are established on one model; the cross-model sweeps are coarse, and on Qwen2.5-7B they do not resolve the read boundary. White-box evidence is at 4B to 14B; frontier evidence is behavioral. The read-versus-recompute results support within-experiment contrasts, and the only patch control is size-matched noise; a structured control planting a valid activation for a different quantity would further isolate value content.

9 Conclusion

Serial reasoning state in language models has no internal home: closing the written channel reveals no fallback capacity at any scale, writing begins at no threshold, and no intervention shifts state between the workspace and the trace. The chain of thought is the working memory this leaves, and it works as compositional registers: settable by intervention, retrieved through attention in the middle layers, independently addressable, and decisive for the answer. For oversight the implication cuts both ways: the written positions are a causally used memory channel, and identical text can carry different stored values, so trace monitoring observes the channel while verifying the state behind it remains a separate problem.

References

- Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. Causal abstractions of neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021. arXiv:2106.02997.
- Wes Gurnee et al. Reading concepts from the residual stream with jacobian lenses. *arXiv preprint*, 2026.
- Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv preprint arXiv:2404.15255*, 2024.
- Jack Lanchantin, Shubham Toshniwal, Jason Weston, Arthur Szlam, and Sainbayar Sukhbaatar. Learning to reason and memorize with self-notes. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.00833.
- Tamera Lanham et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2202.05262.
- William Merrill and Ashish Sabharwal. The expressive power of transformers with chain of thought. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.07923.
- Jacob Pfau, William Merrill, and Samuel R. Bowman. Let’s think dot by dot: Hidden computation in transformer language models. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2404.15758.
- Andy Shih et al. Do models read what they write? causal registers in scratchpad reasoning. *arXiv preprint arXiv:2606.29522*, 2026.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.04388.
- Fangwei Zhu et al. Chain-of-thought tokens are computer program variables. *arXiv preprint arXiv:2505.04955*, 2025.